

Final Project Group 7

Object Detection Using YOLO


Hao Heng, Ze Gong, Junhe Zhang

The bottom right corner of the slide features a series of overlapping, semi-transparent blue geometric shapes, including triangles and parallelograms, creating a modern, abstract design.

Introduction

- YOLO
- Create custom dataset
- Train YOLO network
- Result and Conclusion

What is Yolo?



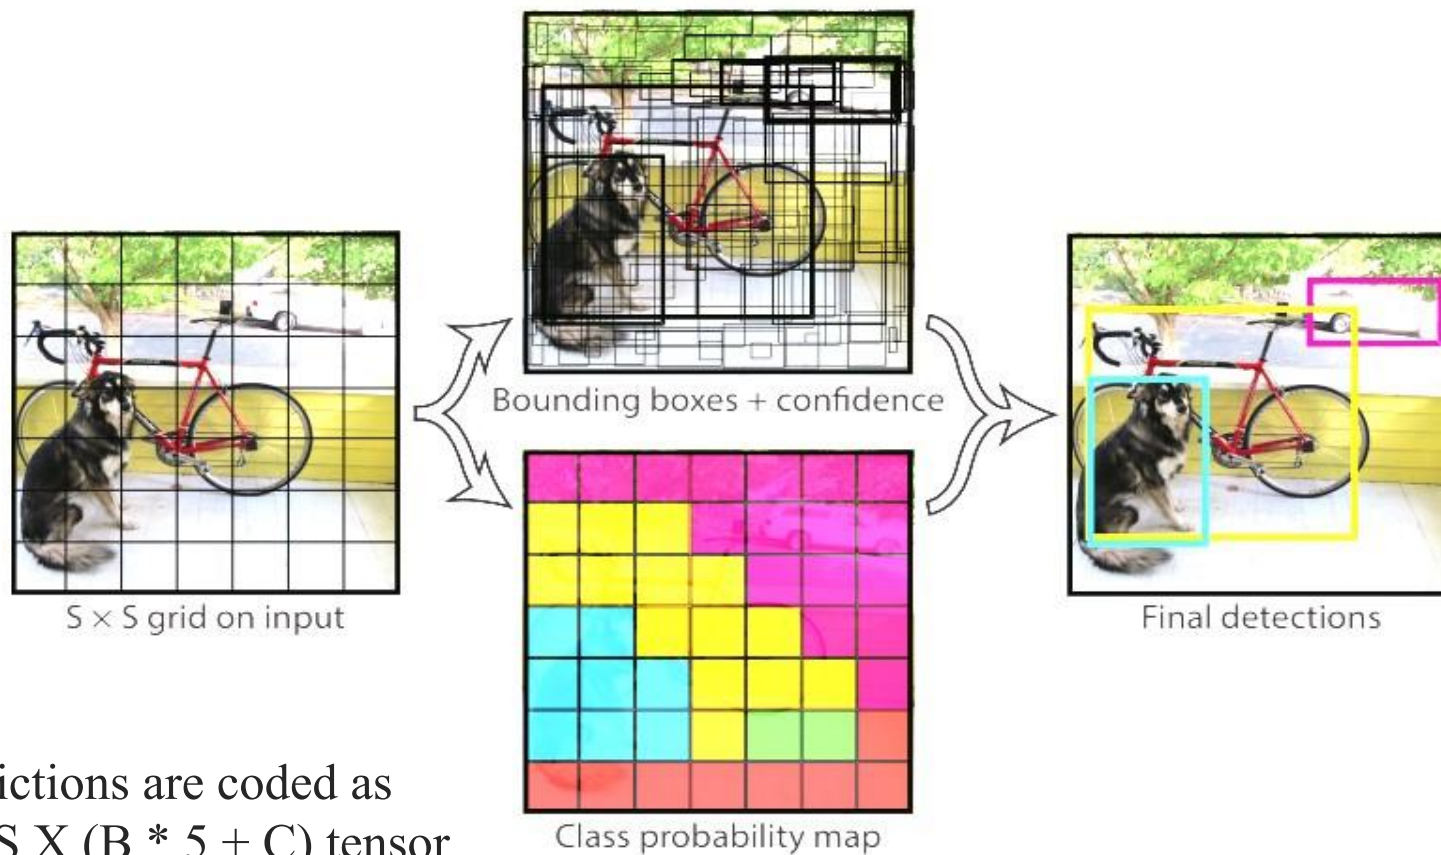
- You only look once (YOLO) is a state-of-the-art, real-time object detection system

YOLO Network


| | Type | Filters | Size | Output |
|----|---------------|---------|------------------|------------------|
| | Convolutional | 32 | 3×3 | 256×256 |
| | Convolutional | 64 | $3 \times 3 / 2$ | 128×128 |
| 1x | Convolutional | 32 | 1×1 | 128×128 |
| | Convolutional | 64 | 3×3 | |
| | Residual | | | |
| | Convolutional | 128 | $3 \times 3 / 2$ | 64×64 |
| 2x | Convolutional | 64 | 1×1 | 64×64 |
| | Convolutional | 128 | 3×3 | |
| | Residual | | | |
| | Convolutional | 256 | $3 \times 3 / 2$ | 32×32 |
| 8x | Convolutional | 128 | 1×1 | 32×32 |
| | Convolutional | 256 | 3×3 | |
| | Residual | | | |
| | Convolutional | 512 | $3 \times 3 / 2$ | 16×16 |
| 8x | Convolutional | 256 | 1×1 | 16×16 |
| | Convolutional | 512 | 3×3 | |
| | Residual | | | |
| | Convolutional | 1024 | $3 \times 3 / 2$ | 8×8 |
| 4x | Convolutional | 512 | 1×1 | 8×8 |
| | Convolutional | 1024 | 3×3 | |
| | Residual | | | |
| | Avgpool | | Global | |
| | Connected | | 1000 | |
| | Softmax | | | |

In this project we use YOLO (You Only Look Once) v3.0 for the purpose of Object Detection.

YOLO



Predictions are coded as
 $S \times S \times (B * 5 + C)$ tensor



Calculation of the bounding box coordinates.

Create custom dataset

Dataset

Source: Google Open Image Dataset

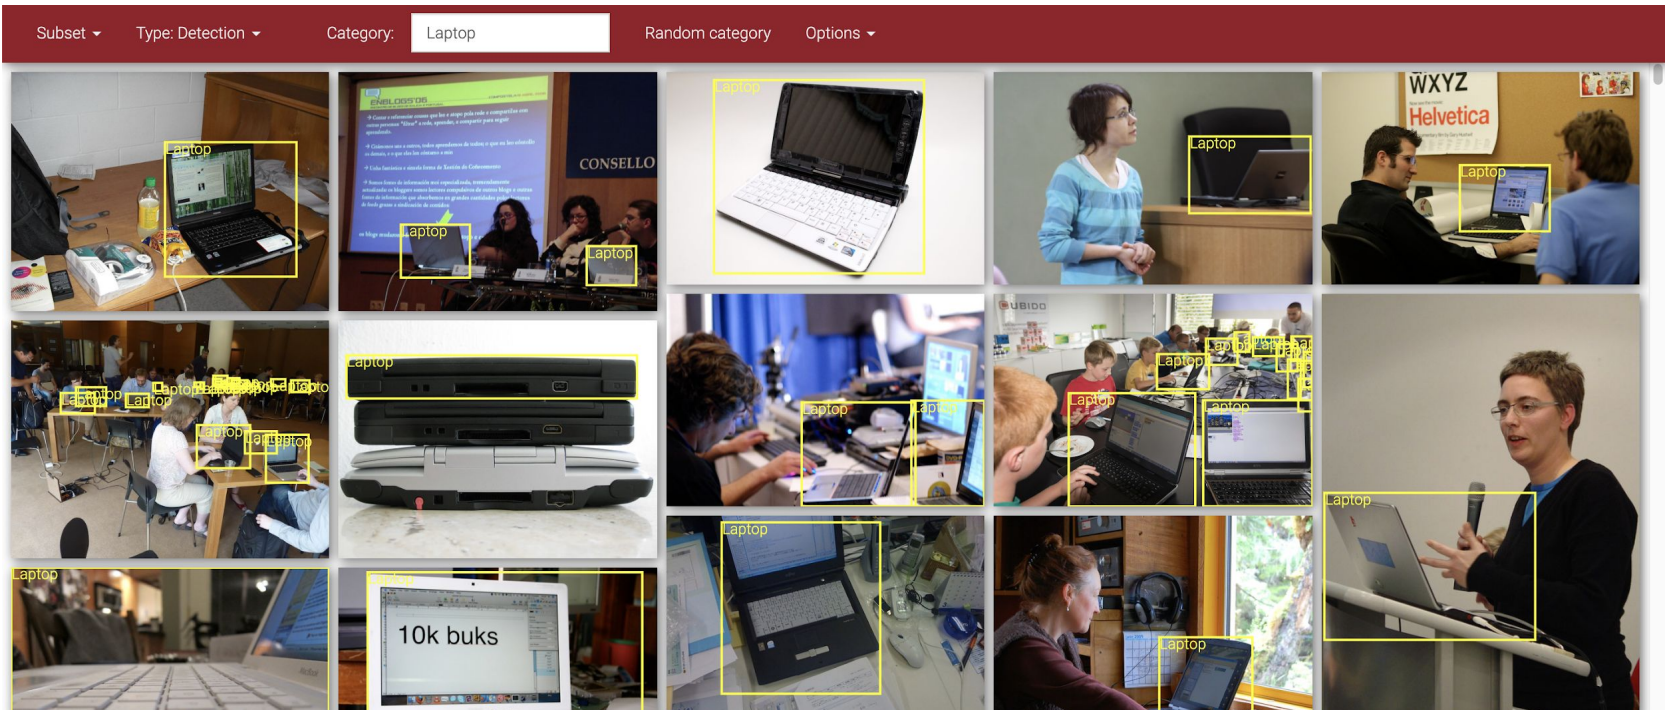
link: <https://storage.googleapis.com/openimages/web/index.html>

Downloader: OIv4_ToolKit

Category: Car, Bicycle wheel, Traffic light, Person, Jeans, Laptop, Bus

Instance of each category: 1000

Dataset



Data preprocessing

Needed files for darknet framework:

- annotation for each image
- train.txt
- test.txt
- data.data
- classes.names

Data preprocessing

train.txt:

/full path/image1.jpg

/full path/image2.jpg

.....

test.txt:

/full path/image1.jpg

/full path/image2.jpg

.....

```
ubuntu@ip-172-31-21-176:~/open-image-data/OIDv4_ToolKit/OID/Dataset/train$ cat train.txt
/home/ubuntu/open-image-data/OIDv4_ToolKit/OID/Dataset/train/Car_Bicycle_wheel_Bus_Traffic_light_Jeans_Laptop_Person/94c83d82c3a4938.jpg
/home/ubuntu/open-image-data/OIDv4_ToolKit/OID/Dataset/train/Car_Bicycle_wheel_Bus_Traffic_light_Jeans_Laptop_Person/8e160ac1db81e696.jpg
/home/ubuntu/open-image-data/OIDv4_ToolKit/OID/Dataset/train/Car_Bicycle_wheel_Bus_Traffic_light_Jeans_Laptop_Person/8ef497da3e77c268.jpg
/home/ubuntu/open-image-data/OIDv4_ToolKit/OID/Dataset/train/Car_Bicycle_wheel_Bus_Traffic_light_Jeans_Laptop_Person/c4fbc9582192f13e.jpg
/home/ubuntu/open-image-data/OIDv4_ToolKit/OID/Dataset/train/Car_Bicycle_wheel_Bus_Traffic_light_Jeans_Laptop_Person/d9cd1988bd1ac21f.jpg
/home/ubuntu/open-image-data/OIDv4_ToolKit/OID/Dataset/train/Car_Bicycle_wheel_Bus_Traffic_light_Jeans_Laptop_Person/f33d1607c9a993c4.jpg
/home/ubuntu/open-image-data/OIDv4_ToolKit/OID/Dataset/train/Car_Bicycle_wheel_Bus_Traffic_light_Jeans_Laptop_Person/eeb824867939c14a.jpg
/home/ubuntu/open-image-data/OIDv4_ToolKit/OID/Dataset/train/Car_Bicycle_wheel_Bus_Traffic_light_Jeans_Laptop_Person/5383c88b95ac77ff.jpg
/home/ubuntu/open-image-data/OIDv4_ToolKit/OID/Dataset/train/Car_Bicycle_wheel_Bus_Traffic_light_Jeans_Laptop_Person/5d254a888ea6196d.jpg
/home/ubuntu/open-image-data/OIDv4_ToolKit/OID/Dataset/train/Car_Bicycle_wheel_Bus_Traffic_light_Jeans_Laptop_Person/7bc5f899de1f888b.jpg
/home/ubuntu/open-image-data/OIDv4_ToolKit/OID/Dataset/train/Car_Bicycle_wheel_Bus_Traffic_light_Jeans_Laptop_Person/a7c2706b8db095778b.jpg
/home/ubuntu/open-image-data/OIDv4_ToolKit/OID/Dataset/train/Car_Bicycle_wheel_Bus_Traffic_light_Jeans_Laptop_Person/a042e475d4429f8a.jpg
/home/ubuntu/open-image-data/OIDv4_ToolKit/OID/Dataset/train/Car_Bicycle_wheel_Bus_Traffic_light_Jeans_Laptop_Person/0da97a61df1b8a1e.jpg
/home/ubuntu/open-image-data/OIDv4_ToolKit/OID/Dataset/train/Car_Bicycle_wheel_Bus_Traffic_light_Jeans_Laptop_Person/fe54cc512921db99.jpg
/home/ubuntu/open-image-data/OIDv4_ToolKit/OID/Dataset/train/Car_Bicycle_wheel_Bus_Traffic_light_Jeans_Laptop_Person/f201af7676a4e7d7.jpg
/home/ubuntu/open-image-data/OIDv4_ToolKit/OID/Dataset/train/Car_Bicycle_wheel_Bus_Traffic_light_Jeans_Laptop_Person/4aa753a7ba32ec9.jpg
/home/ubuntu/open-image-data/OIDv4_ToolKit/OID/Dataset/train/Car_Bicycle_wheel_Bus_Traffic_light_Jeans_Laptop_Person/35672503659685a7.jpg
/home/ubuntu/open-image-data/OIDv4_ToolKit/OID/Dataset/train/Car_Bicycle_wheel_Bus_Traffic_light_Jeans_Laptop_Person/9c836ecb6a31ee62.jpg
/home/ubuntu/open-image-data/OIDv4_ToolKit/OID/Dataset/train/Car_Bicycle_wheel_Bus_Traffic_light_Jeans_Laptop_Person/62210e2a47952656.jpg
```

Data preprocessing

data.data:

classes=the number of classes in the dataset

train=full path to train.txt file

valid=full path to test.txt file

```
Personubuntu@ip-172-31-21-176:~/open-image-data/OIDv4_ToolKit/OID/Dataset/train/Car_Bicycle_wheel_Bus_Traffic_light_Jeans_Laptop_Person$ cat custom_data.data
classes = 7
train = /home/ubuntu/open-image-data/OIDv4_ToolKit/OID/Dataset/train/Car_Bicycle_wheel_Bus_Traffic_light_Jeans_Laptop_Person/train.txt
valid = /home/ubuntu/open-image-data/OIDv4_ToolKit/OID/Dataset/train/Car_Bicycle_wheel_Bus_Traffic_light_Jeans_Laptop_Person/test.txt
names = /home/ubuntu/open-image-data/OIDv4_ToolKit/OID/Dataset/train/Car_Bicycle_wheel_Bus_Traffic_light_Jeans_Laptop_Person/classes.names
backup = backupubuntu@ip-172-31-21-176:~/open-image-data/OIDv4_ToolKit/OID/Dataset/train/Car_Bicycle_wheel_Bus_Traffic_light_Jeans_Laptop_Person$
```

classes.names:

Car

Bicycle wheel

Data preprocessing

```
Personubuntu@ip-172-31-21-176:~/open-image-data/OIDv4_ToolKit/OID/Dataset/train/Car_Bicycle_wheel_Bus_Traffic_light_Jeans_Laptop_Person$ cat custom_data.data
classes = 7
train = /home/ubuntu/open-image-data/OIDv4_ToolKit/OID/Dataset/train/Car_Bicycle_wheel_Bus_Traffic_light_Jeans_Laptop_Person/train.txt
valid = /home/ubuntu/open-image-data/OIDv4_ToolKit/OID/Dataset/train/Car_Bicycle_wheel_Bus_Traffic_light_Jeans_Laptop_Person/test.txt
names = /home/ubuntu/open-image-data/OIDv4_ToolKit/OID/Dataset/train/Car_Bicycle_wheel_Bus_Traffic_light_Jeans_Laptop_Person/classes.names
backup = backupubuntu@ip-172-31-21-176:~/open-image-data/OIDv4_ToolKit/OID/Dataset/train/Car_Bicycle_wheel_Bus_Traffic_light_Jeans_Laptop_Person$
```

```
./home, ubuntu, open-image-data
Car
Bicycle wheel
Bus
Traffic light
Jeans
Laptop
Personubuntu@ip-172-31-21-176:~/open-image-data/OIDv4_ToolKit/OID/Dataset/train/Car_Bicycle_wheel_Bus_Traffic_light_Jeans_Laptop_Person$
```

Training YOLO in Darknet Framework

Setting up Configuration

Classes = 7

filters = (7 + 5) x 3 = 36

max_batches = 7 x 2000 = 14000

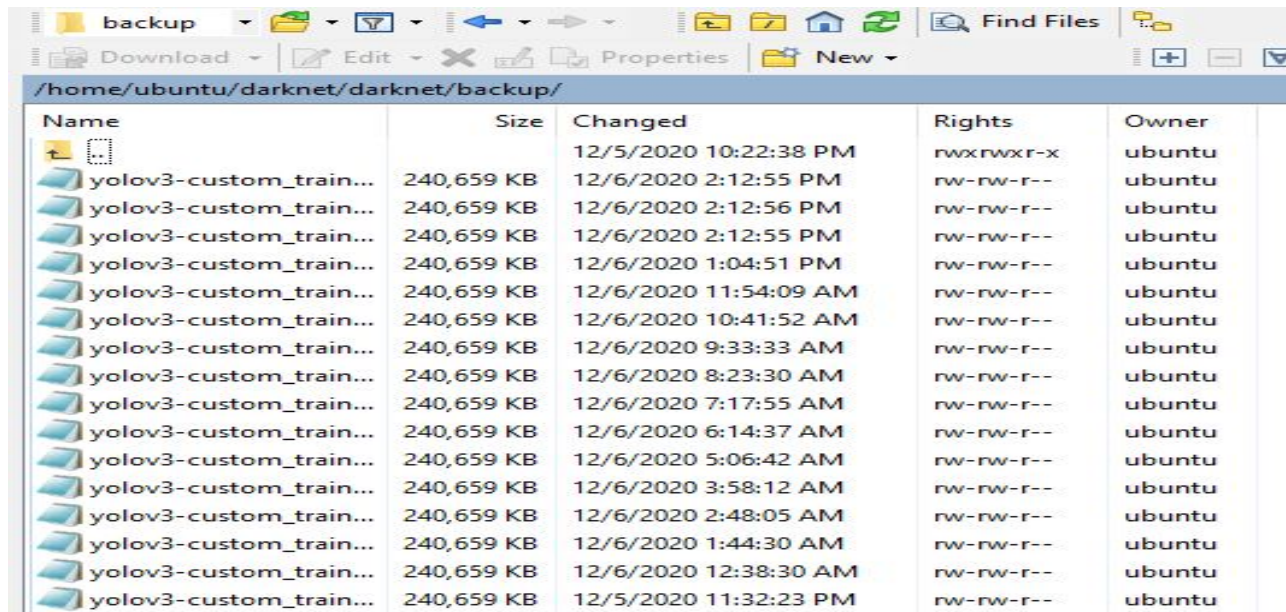
steps = 0.8 x max_batches / 0.9 x max_batches
= 11200 / 12600

batch = 32

subdivisions = 8

minibatch = batch/subdivisions = 4

Result of Training Process



The screenshot shows a file manager window titled 'backup' with the address bar displaying '/home/ubuntu/darknet/darknet/backup/'. The window contains a table of files, all named 'yolov3-custom_train...', each with a size of 240,659 KB. The files are sorted by modification time, ranging from 12/5/2020 10:22:38 PM to 12/5/2020 11:32:23 PM. All files have 'rw-rw-r--' permissions and are owned by 'ubuntu'.

| Name | Size | Changed | Rights | Owner |
|------------------------|------------|-----------------------|-----------|--------|
| ↑ | | 12/5/2020 10:22:38 PM | rw-rw-r-- | ubuntu |
| yolov3-custom_train... | 240,659 KB | 12/6/2020 2:12:55 PM | rw-rw-r-- | ubuntu |
| yolov3-custom_train... | 240,659 KB | 12/6/2020 2:12:56 PM | rw-rw-r-- | ubuntu |
| yolov3-custom_train... | 240,659 KB | 12/6/2020 2:12:55 PM | rw-rw-r-- | ubuntu |
| yolov3-custom_train... | 240,659 KB | 12/6/2020 1:04:51 PM | rw-rw-r-- | ubuntu |
| yolov3-custom_train... | 240,659 KB | 12/6/2020 11:54:09 AM | rw-rw-r-- | ubuntu |
| yolov3-custom_train... | 240,659 KB | 12/6/2020 10:41:52 AM | rw-rw-r-- | ubuntu |
| yolov3-custom_train... | 240,659 KB | 12/6/2020 9:33:33 AM | rw-rw-r-- | ubuntu |
| yolov3-custom_train... | 240,659 KB | 12/6/2020 8:23:30 AM | rw-rw-r-- | ubuntu |
| yolov3-custom_train... | 240,659 KB | 12/6/2020 7:17:55 AM | rw-rw-r-- | ubuntu |
| yolov3-custom_train... | 240,659 KB | 12/6/2020 6:14:37 AM | rw-rw-r-- | ubuntu |
| yolov3-custom_train... | 240,659 KB | 12/6/2020 5:06:42 AM | rw-rw-r-- | ubuntu |
| yolov3-custom_train... | 240,659 KB | 12/6/2020 3:58:12 AM | rw-rw-r-- | ubuntu |
| yolov3-custom_train... | 240,659 KB | 12/6/2020 2:48:05 AM | rw-rw-r-- | ubuntu |
| yolov3-custom_train... | 240,659 KB | 12/6/2020 1:44:30 AM | rw-rw-r-- | ubuntu |
| yolov3-custom_train... | 240,659 KB | 12/6/2020 12:38:30 AM | rw-rw-r-- | ubuntu |
| yolov3-custom_train... | 240,659 KB | 12/5/2020 11:32:23 PM | rw-rw-r-- | ubuntu |

Mean Average Precision

Example:

| | |
|---|----------------------|
| class_id = 0, name = Car, ap = 51.01% | (TP = 308, FP = 165) |
| class_id = 1, name = Bicycle wheel, ap = 62.46% | (TP = 304, FP = 68) |
| class_id = 2, name = Bus, ap = 79.25% | (TP = 161, FP = 35) |
| class_id = 3, name = Traffic light, ap = 44.14% | (TP = 329, FP = 222) |
| class_id = 4, name = Jeans, ap = 55.24% | (TP = 183, FP = 64) |
| class_id = 5, name = Laptop, ap = 80.80% | (TP = 155, FP = 28) |
| class_id = 6, name = Person, ap = 23.76% | (TP = 230, FP = 308) |


for conf_thresh = 0.25, precision = 0.65, recall = 0.45, F1-score = 0.53

for conf_thresh = 0.25, TP = 1670, FP = 890, FN = 2047, average IoU = 50.15 %

IoU threshold = 50 %, used Area-Under-Curve for each unique Recall

mean average precision (mAP@0.50) = 0.566669, or 56.67 %

Performance of Different Training Weights



The best mAP is 56.82%, appears at iteration 12000.

Apply The Best Model on Testing image and video

Loading Best Model

There are two different ways to do object detection:

- Using darknet framework's command line
 - `./darknet detector test cfg/custom_data.data cfg/yolov3-custom_train.cfg backup/yolov3-custom_train_last.weights data/laptop-jean-test.jpg -out_filename data/result-laptop-jean-test.jpg -dont_show`
- Run `process_image_video.py` file
 - `python3 process_image_video.py`

Algorithm of process_image_video.py

1. Read the input image and convert it into a blob object which is accepted by the YOLO framework.
2. Load the best trained weights and the network structure file into the YOLO framework.
3. Process blob image file into the model and generate output bounding boxes for each object inside the image.
4. Using Non-Maximun_Suppression technique on the output bounding boxes to drop duplicated boxes on the same object with lower confidence.



What the Results look like

```
out_scores, out_boxes, out_classes = predict(sess, "test.jpg")
```

```
Found 7 boxes for test.jpg  
car 0.60 (925, 285) (1045, 374)  
car 0.66 (706, 279) (786, 350)  
bus 0.67 (5, 266) (220, 407)  
car 0.70 (947, 324) (1280, 705)  
car 0.74 (159, 303) (346, 440)  
car 0.80 (761, 282) (942, 412)  
car 0.89 (367, 300) (745, 648)
```



Non-Maximum-Suppression



Intersection over Union

$$IoU = \frac{B_1 \cap B_2}{B_1 \cup B_2} = \frac{\text{Intersection}}{\text{Union}}$$

Results of Object Detected Images




Original Image



Model Detection

Results of Object Detected Images

Original Image



Model Detection

Results of Object Detected Images



Original Image



Model Detection

Results of Object Detected Videos



<https://drive.google.com/file/d/1luAp5ggEG8hvPBIzG8yjvvBXh1PIXCIW/view?usp=sharing>



<https://drive.google.com/file/d/1xc2RdESogHPag3VwwnbuNQs-RX--mZzF/view?usp=sharing>

Limitation and Conclusion

Limitations to YOLO

In spite of the high accuracy numbers and almost perfect bounding boxes, there are a few limitations to this object detection methods:

- Sometimes fails to detect overlapping objects, and objects that are partially visible in the frame.
- Detects the object class falsely if the features are a little blurred.
- It is very sensitive to model overfitting.

Conclusion and discussion

- The number of iteration of training process is not long enough to predict all categories we choose. (2000 iteration for each class)
- Our model is under-trained, compared to coco weight. Our dataset is not big enough to train a very accurate model
- Our model performs not well at detecting bus and person
- Our model has some difficult to detect object when there are multiple classes are overlapped. It will only detect one class.

References

- [1] Joseph Redmon, Santosh Divvala, Ross Girshick, Ali Farhadi “You Only Look Once: Unified, Real-Time Object Detection”, IEEE Conference on Computer Vision and Pattern Recognition (CVPR), 2016.

- [2] Open Images Dataset V6 + Extensions. (n.d.). Retrieved December 08, 2020, from <https://storage.googleapis.com/openimages/web/index.html>

- [3] Joseph R. . Darknet: Open Source Neural Networks in C. Retrieved (2013-2016) from <http://pjreddie.com/darknet/>

- [4] Valentyn S. Training YOLO v3 for Objects Detection with Custom Data.
Retrieved from <https://storage.googleapis.com/openimages/web/index.html>

Thank you!